

Varnish 2.0.x and eZ publish 4.0.x

Technical implementation guide for running Varnish and eZ publish on the same server to achieve high performance on your eZ publish web site.

Server-environment for this use-case is Debian Linux Etch with the deb-packages for apache2, php5, mysql-server 5 and related packages. Varnish is compiled from source because the deb-package of Varnish is too old.

Ref: <http://varnish.projects.linpro.no/>, and <http://ez.no/ezpublish>

«Varnish is a state-of-the-art, high-performance HTTP accelerator. It uses the advanced features in Linux 2.6, FreeBSD 6/7 and Solaris 10 to achieve its high performance.»

«eZ Publish is an Open Source Content Management System chosen by [thousands of enterprises and organizations](#) world wide. It helps you build corporate websites, [intranets](#), [webshops](#) and [media portals](#). eZ Publish is 100% Open Source, available either as a [free download](#) or as an enterprise solution [“eZ Publish Premium”](#) with support, guarantees and maintenance.»

My example doesn't utilize cache-purging on publish. I have configured Varnish to automatically purge cache for content after 5 minutes, and images, media, css and js are purges every 10 minutes.

Purge on publish should be possible to setup by integrating the eZ extension:

<http://projects.ez.no/all2evcc>, but I have not done it yet.

Varnish installation

Grab Varnish 2.0.x from source and compile and install following the Varnish Installation documentation at <http://varnish.projects.linpro.no/wiki/Installation>

Create configuration file /etc/default/varnish and /etc/varnish/vcl.conf

Create startup-file /etc/init.d/varnish and symlink to startup-file in default runlevel to start varnish on server boot.

Varnish configuration

Example configuration file /etc/default/varnish:

```
# Configuration file for varnish
#
# /etc/init.d/varnish expects the variable $DAEMON_OPTS to be set from this
# shell script fragment.
#
# Maximum number of open files (for ulimit -n)
NFILES=131072
```

```
# Default varnish instance name is the local nodename. Can be overridden with
# the -n switch, to have more instances on a single server.
INSTANCE=$(uname -n)

# This file contains 4 alternatives, please use only one.

## Alternative 1, Minimal configuration, no VCL
#
# Listen on port 6081, administration on localhost:6082, and forward to
# content server on localhost:8080. Use a fixed-size cache file.
#
#DAEMON_OPTS="-a :80 \
#             -T localhost:6082 \
#             -b localhost:81 \
#             -u www-data -g www-data \
#             -s file,/var/lib/varnish/$INSTANCE/varnish_storage.bin,1G"

## Alternative 2, Configuration with VCL
#
# Listen on port 6081, administration on localhost:6082, and forward to
# one content server selected by the vcl file, based on the request. Use a
# fixed-size cache file.
#
DAEMON_OPTS="-a :80 \
             -T localhost:6082 \
             -f /etc/varnish/vcl.conf \
             -s file,/var/lib/varnish/$INSTANCE/varnish_storage.bin,1G"

## Alternative 3, Advanced configuration
#
# See varnishd(1) for more information.
#
# # Main configuration file. You probably want to change it :)
# VARNISH_VCL_CONF=/etc/varnish/default.vcl
#
# # Default address and port to bind to
# # Blank address means all IPv4 and IPv6 interfaces, otherwise specify
# # a host name, an IPv4 dotted quad, or an IPv6 address in brackets.
```

```
# VARNISH_LISTEN_ADDRESS=
# VARNISH_LISTEN_PORT=6081
#
# # Telnet admin interface listen address and port
# VARNISH_ADMIN_LISTEN_ADDRESS=127.0.0.1
# VARNISH_ADMIN_LISTEN_PORT=6082
#
# # The minimum number of worker threads to start
# VARNISH_MIN_THREADS=1
#
# # The Maximum number of worker threads to start
# VARNISH_MAX_THREADS=1000
#
# # Idle timeout for worker threads
# VARNISH_THREAD_TIMEOUT=120
#
# # Cache file location
# VARNISH_STORAGE_FILE=/var/lib/varnish/$INSTANCE/varnish_storage.bin
#
# # Cache file size: in bytes, optionally using k / M / G / T suffix,
# # or in percentage of available disk space using the % suffix.
# VARNISH_STORAGE_SIZE=1G
#
# # Backend storage specification
# VARNISH_STORAGE="file,${VARNISH_STORAGE_FILE},${VARNISH_STORAGE_SIZE}"
#
# # Default TTL used when the backend does not specify one
# VARNISH_TTL=120
#
# # DAEMON_OPTS is used by the init script.  If you add or remove options, make
# # sure you update this section, too.
# DAEMON_OPTS="-a ${VARNISH_LISTEN_ADDRESS}:${VARNISH_LISTEN_PORT} \
#             -f ${VARNISH_VCL_CONF} \
#             -T ${VARNISH_ADMIN_LISTEN_ADDRESS}:${VARNISH_ADMIN_LISTEN_PORT} \
#             -t ${VARNISH_TTL} \
#             -w ${VARNISH_MIN_THREADS},${VARNISH_MAX_THREADS},${VARNISH_THREAD_TIMEOUT} \
#             -s ${VARNISH_STORAGE}"
#
## Alternative 4, Do It Yourself
```

```
#  
# DAEMON_OPTS=""
```

Example configuration file /etc/varnish/vcl.conf:

```
# Default backend definition.  Set this to point to your content  
# server.  
  
backend b1 {  
    .host = "127.0.0.1";  
    .port = "6081";  
}  
  
acl purge {  
    "localhost";  
}  
  
## Called when a client request is received  
#  
#pipe = just a pipe through varnish. No checks whatsoever  
#pass = Go through all checks -> everything, just dont lookup in the cache  
  
sub vcl_recv {  
  
    # Add a unique header containing the client address  
    unset req.http.X-Forwarded-For;  
    set    req.http.X-Forwarded-For = client.ip;  
  
#### always cache these items:  
  
    ## images  
    if (req.request == "GET" && req.url ~ "\.(gif|jpg|jpeg|bmp|png|tiff|tif|ico|img|  
tga|wmf)$") {  
        set req.backend = b1;  
        lookup;  
    }  
}
```

```
## various other content pages
if (req.request == "GET" && req.url ~ "\.(css|js)$") {
    set req.backend = b1;
    lookup;
}

## multimedia
if (req.request == "GET" && req.url ~ "\.(svg|swf|mov|avi|wmv)$") {
    set req.backend = b1;
    lookup;
}

#### do not cache these files

#     if (req.request == "GET" && req.url ~ "\.(cfm)$") {
#         pass;
#     }

#### do not cache these rules:

# pass mode can't handle POST (yet)
if (req.request == "POST") {
    pipe;
}

if (req.request != "GET" && req.request != "HEAD") {
    pipe;
}

if (req.http.Expect) {
    pipe;
}

if (req.http.Authenticate || req.http.Authorization) {
    pass;
}

#### if there is a purge make sure its coming from $localhost

if (req.request == "PURGE") {
    if(!client.ip ~ purge) {
        error 405 "Not Allowed";
    }
}
```

```
        }
        purge_url(req.http.X-Purge-Url);
        error 200 "Purged";
    }

####

#### unknown function to "normalize the Accept-Encoding headers"

    if (req.http.Accept-Encoding) {
        if (req.http.Accept-Encoding ~ "gzip") {
            set req.http.Accept-Encoding = "gzip";
        } elseif (req.http.Accept-Encoding ~ "deflate") {
            set req.http.Accept-Encoding = "deflate";
        } else {
            # unkown algorithm
            unset req.http.Accept-Encoding;
        }
    }

#### don't cache authenticated sessions
    if (req.http.Cookie && req.http.Cookie ~ "is_logged_in=") {
        pipe;
    }

    // Varnish doesn't do INM requests so pass it through if no If-Modified-Since
was sent
    if (req.http.If-None-Match && !req.http.If-Modified-Since) {
        pass;
    }

#### if it passes all these tests, do a lookup anyway;
    lookup;

}

#
## Called when entering pipe mode
#
sub vcl_pipe {
    pipe;
}
```

```
#
## Called when entering pass mode
#
sub vcl_pass {
    pass;
}

#
## Called when entering an object into the cache
#
sub vcl_hash {
    set req.hash += req.url;
    set req.hash += req.http.host;
    hash;
}

#
## Called when the requested object was found in the cache
#
sub vcl_hit {

    # if (obj.http.X-Cache == "MISS") {
    #     set obj.http.X-Cache = "HIT";
    # }

    if (!obj.cacheable) {
        pass;
    }

    deliver;
}

#
## Called when the requested object was not found in the cache
#
sub vcl_miss {
    fetch;
}

#
## Called when the requested object has been retrieved from the
## backend, or the request to the backend has failed
```

```
#
sub vcl_fetch {
    # default time to live for cache objects
    set obj.ttl = 300s;

    if (!obj.cacheable) {
        pass;
    }

    if (obj.http.Set-Cookie ~ "is_logged_in=deleted(.*)" ) {
        deliver;
    }

    if (obj.http.Set-Cookie) {
        pass;
    }

    if (req.request == "GET" && req.url ~ "\.(gif|jpg|jpeg|bmp|png|tiff|tif|ico|img|tga|wmf)
    $") {
        set obj.ttl = 600s;
        deliver;
    }

    ## various other content pages
    if (req.request == "GET" && req.url ~ "\.(css|js|html)$") {
set obj.ttl = 600s;
        deliver;
    }

    ## multimedia
    if (req.request == "GET" && req.url ~ "\.(svg|swf|ico|mp3|mp4|m4a|ogg|mov|avi|
    wmv)$") {
set obj.ttl = 600s;
        deliver;
    }

    # for varnish 2.0:
    # set obj.prefetch = -30s;

    deliver;
}
```



```
#
#
## Called before a cached object is delivered to the client
#
sub vcl_deliver {
    deliver;
}

#
## Called when an object nears its expiry time
# not supportet yet, see http://varnish.projects.linpro.no/ticket/367
#sub vcl_timeout {
#    fetch;
#}

#
## Called when an object is about to be discarded
#
sub vcl_discard {
    discard;
}
```

Example /etc/init.d/varnish start/stop file:

```
#!/bin/sh
#
# varnish          Control the varnish HTTP accelerator
#
### BEGIN INIT INFO
# Provides:        varnish
# Required-Start:   $local_fs $network
# Required-Stop:    $local_fs $network
# Default-Start:    2 3 4 5
# Default-Stop:     0 1 6
# Short-Description: Start HTTP accelerator
# Description:      This script provides a server-side cache
#                   to be run in front of a httpd and should
#                   listen on port 80 on a properly configured
#                   system
### END INIT INFO
```

```
# Source function library
. /lib/lsb/init-functions

NAME=varnishd
DESC="HTTP accelerator"
PATH=/sbin:/bin:/usr/sbin:/usr/bin:/usr/local/sbin
DAEMON=/usr/local/sbin/varnishd
PIDFILE=/var/run/$NAME.pid

test -x $DAEMON || exit 0

# Include varnish defaults if available
if [ -f /etc/default/varnish ] ; then
    . /etc/default/varnish
fi

# Open files (usually 1024, which is way too small for varnish)
ulimit -n ${NFILES:-131072}

# If $DAEMON_OPTS is not set at all in /etc/default/varnish, use minimal useful
# defaults (Backend at localhost:8080, a common place to put a locally
# installed application server.)
DAEMON_OPTS=${DAEMON_OPTS:--b localhost}

case "$1" in
    start)
        output=$(/bin/tempfile -s.varnish)
        log_daemon_msg "Starting $DESC"
        log_progress_msg $NAME
        if start-stop-daemon \
            --start --quiet --pidfile ${PIDFILE} --exec ${DAEMON} -- \
            -P ${PIDFILE} ${DAEMON_OPTS} > ${output} 2>&1; then
            log_end_msg 0
        else
            log_end_msg 1
            cat $output
            exit 1
        fi
        rm $output
    ;;
    stop)
```

```
        log_daemon_msg "Stopping $DESC"
        log_progress_msg $NAME
        if start-stop-daemon \
            --stop --quiet --pidfile $PIDFILE --retry 10 \
            --exec $DAEMON; then
            log_end_msg 0
        else
            log_end_msg 1
        fi
        ;;
restart|force-reload)
    $0 stop
    $0 start
    ;;
*)
    log_success_msg "Usage: $0 {start|stop|restart|force-reload}"
    exit 1
    ;;
esac

exit 0
```

Apache configuration

To be able to log the actual users in `apachelog` a special logformat needs to be created in `apache`. `Apache` is set up to listen on port 6081, which `Varnish` makes the backend requests to.

Example config file `/etc/apache2/conf.d/apache2.append.conf`

```
LogFormat "%{X-Forwarded-For}i %l %u %t \"%r\" %>s %b \"%{Referer}i\" \"%{User-Agent}i\""
varnishcombined
Listen 6081
NameVirtualHost 127.0.0.1:6081
```

Example virtualhost config file `/etc/apache/sites-available/example.com`

```
<VirtualHost 127.0.0.1:6081>
    KeepAlive Off
    CustomLog /var/log/apache2/www.hit.no.access.log varnishcombined
    ServerName example.com
    DocumentRoot /example.com
    # ...more vhost config directives ...
</VirtualHost>
```

eZ publish configuration

Some configuration of eZ publish must be done to make it behave the right way together with Varnish. Example.com uses frontend editing, and a cookie «is_logged_in» is created by eZ publish and checked by Varnish to identify if a user/editor is logged in. Logged-in sessions are not cached by Varnish (see /etc/varnish/vcl.conf).

Patched ezpublish/index.php with pubsvn.ez.no nextgen rev. 21377

(this will be included standard in ez publish 4.1)

```
        if ( isset( $moduleResult["external_css"] ) )
            $use_external_css = $moduleResult["external_css"];
    }
}

// added rev. 21377 from nextgen: trunk/index.php
$currentUser = eZUser::currentUser();
$ini = eZINI::instance();

$wwwDir = eZSys::wwwDir();
// On host based site accesses this can be empty, causing the cookie to be set
for the current dir,
// but we want it to be set for the whole eZ publish site
$cookiePath = $wwwDir != '' ? $wwwDir : '/';

if ( $currentUser->isLoggedIn() )
{
    setcookie( 'is_logged_in', 'true', 0, $cookiePath );
    header( 'Etag: ' . $currentUser->attribute( 'contentobject_id' ) );
}
else
{
    setcookie( 'is_logged_in', false, 0, $cookiePath );
}

// END rev. 21377.

if ( $module->exitStatus() == eZModule::STATUS_REDIRECT )
```

Example settings/override/site.ini.append.php:

```
[HTTPHeaderSettings]
CustomHeader=enabled
# Cache-Control values are set directly
Cache-Control[]
```

```
Cache-Control[]=
# Expires specifies time offset compared to current time
# Default expired 2 hours ago ( no caching )
Expires[]
Expires[/]=300
# Pragma values are set directly
Pragma[]
Pragma[/]=
```

Example pagelayout.tpl:

```
<html><head>
<meta http-equiv="Pragma" content="no-cache" />
<meta http-equiv="Cache-control" content="no-cache" />
...
</head><body>
...
</body></html>
```